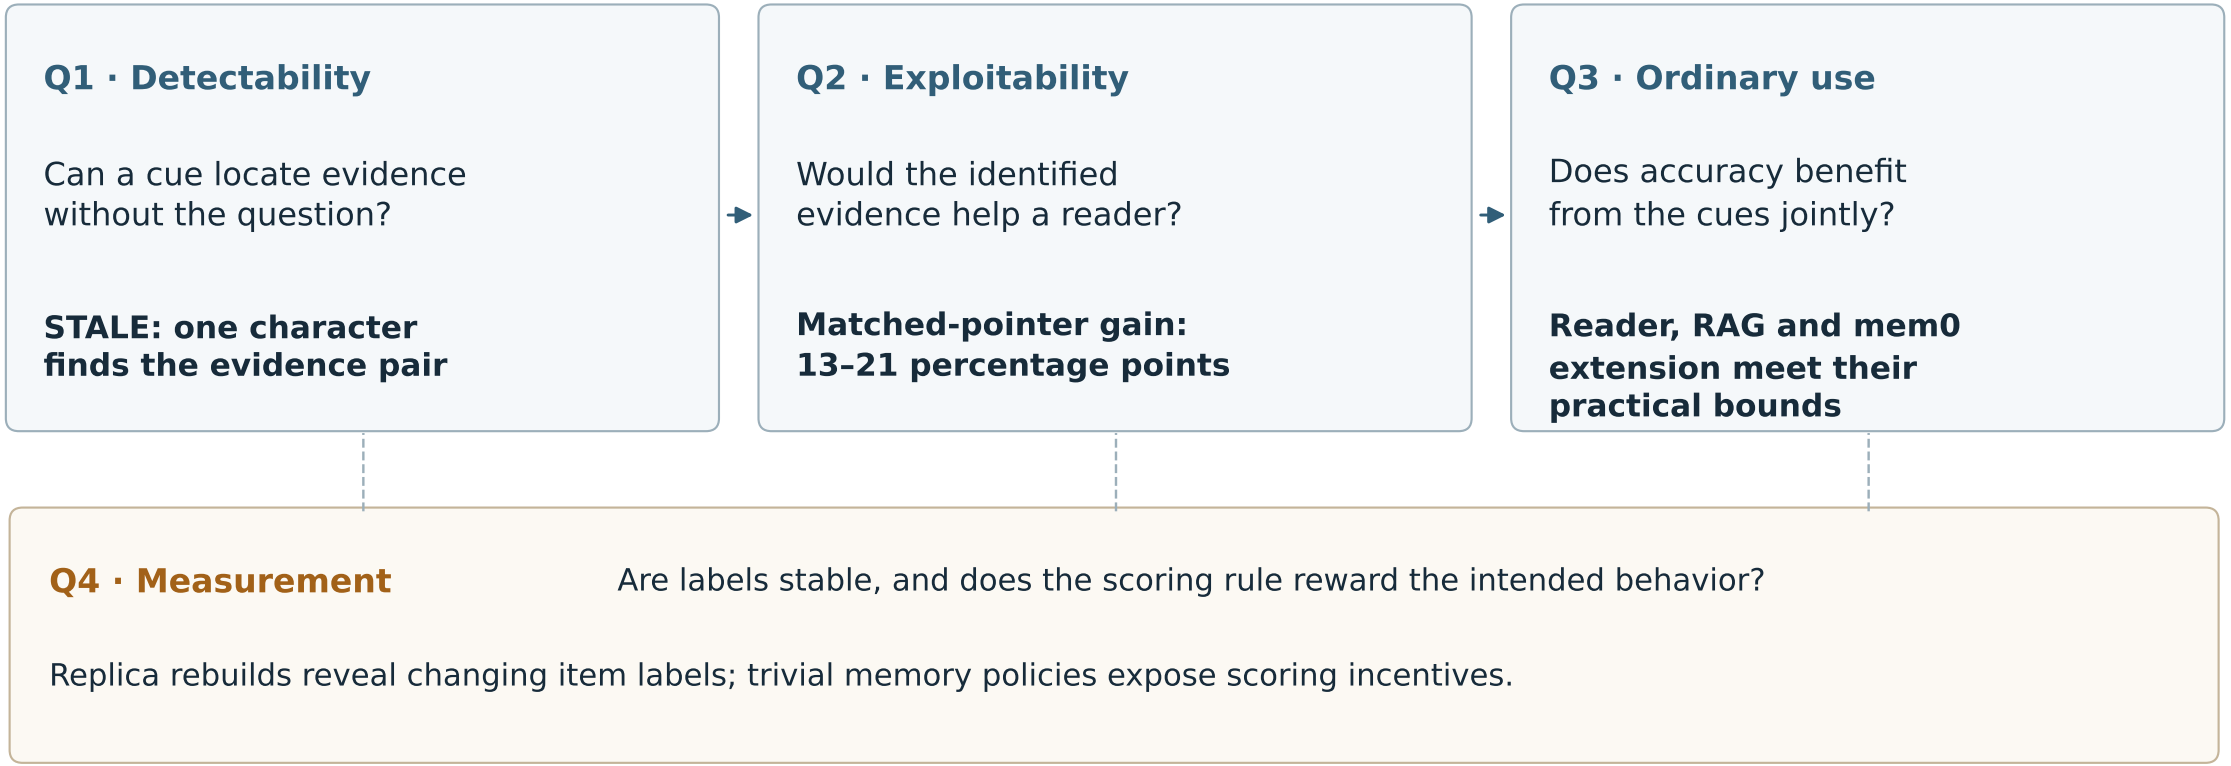
Q1 · Detectability
Can a cue locate evidence without the question?
STALE: one character finds the evidence pair
Q2 · Exploitability
Would the identified evidence help a reader?
Matched-pointer gain: 13–21 percentage points
Q3 · Ordinary use
Does accuracy benefit from the cues jointly?
Reader, RAG and mem0 extension meet their practical bounds
Q4 · Measurement
Are labels stable, and does the scoring rule reward the intended behavior?
Replica rebuilds reveal changing item labels; trivial memory policies expose scoring incentives.