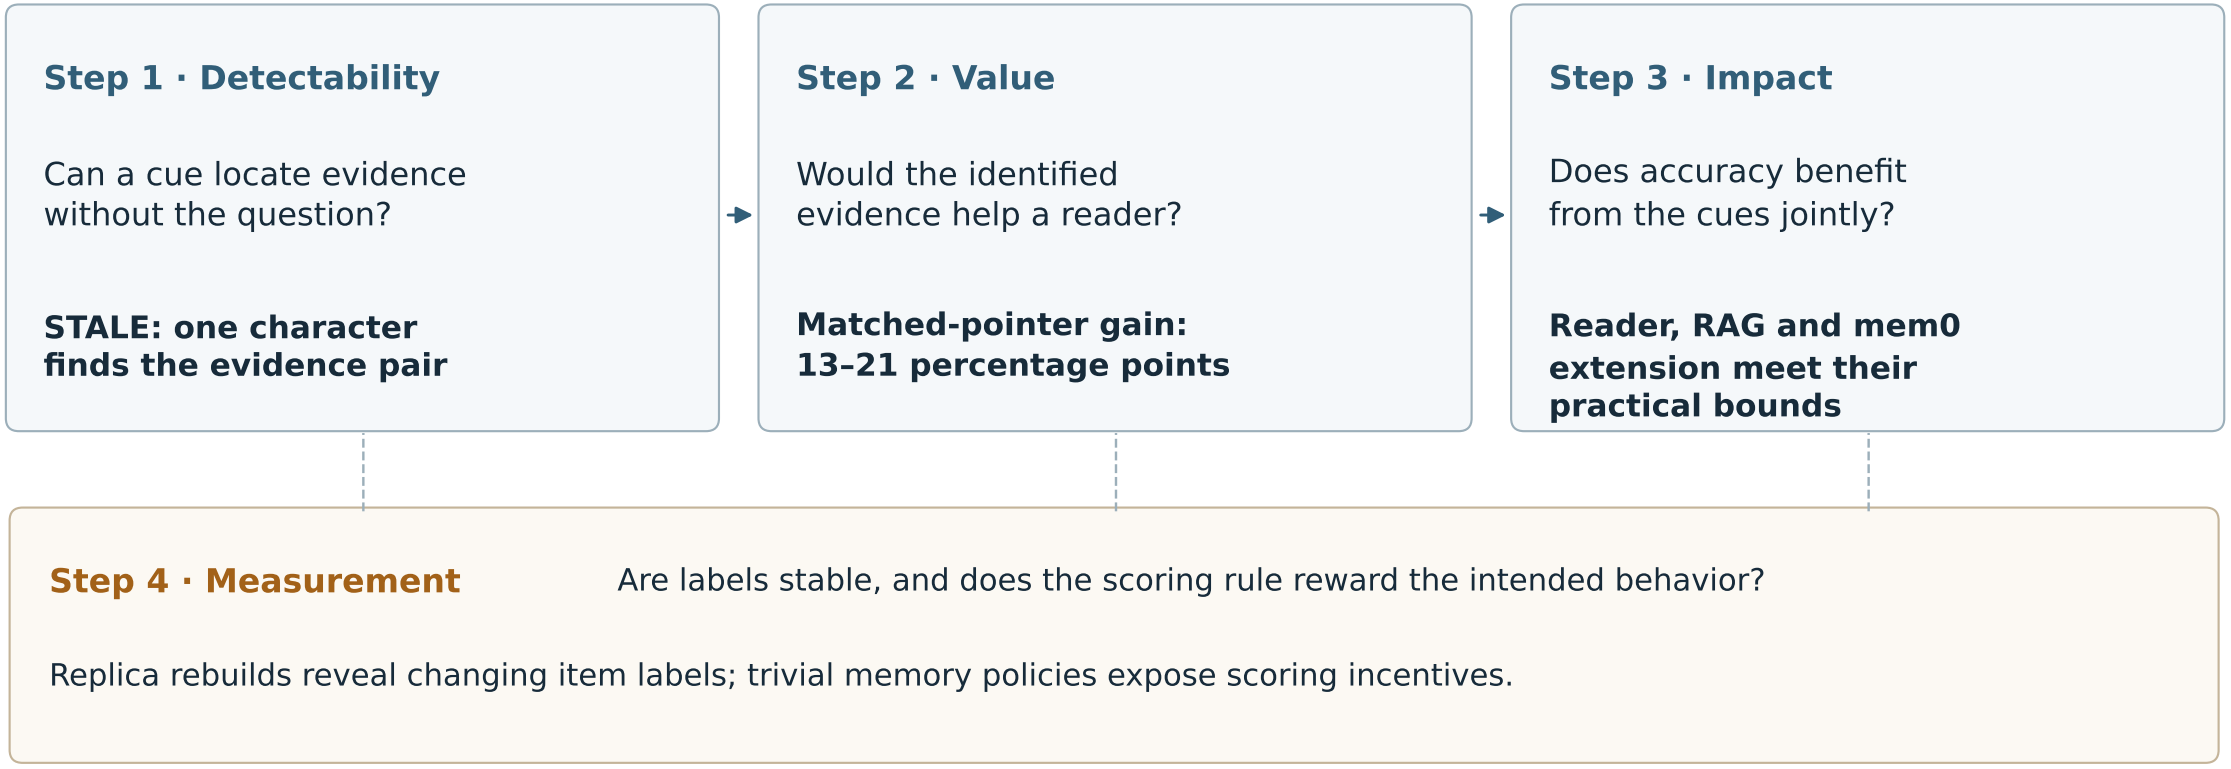
Step 1 · Detectability
Can a cue locate evidence without the question?
STALE: one character finds the evidence pair
Step 2 · Value
Would the identified evidence help a reader?
Matched-pointer gain: 13–21 percentage points
Step 3 · Impact
Does accuracy benefit from the cues jointly?
Reader, RAG and mem0 extension meet their practical bounds
Step 4 · Measurement
Are labels stable, and does the scoring rule reward the intended behavior?
Replica rebuilds reveal changing item labels; trivial memory policies expose scoring incentives.